

Lecture 3 – Coverage Criteria

AAA705: Software Testing and Quality Assurance

Jihyeok Park



2026 Fall

- Random Testing (RT)
 - Geometric distribution – $E(X) = 1/p$, $1 - (1 - p)^n$
 - Residual risk – n clean runs only give $p < \ln(1/\delta)/n$
- Feedback-Directed Random Testing (Randoop)
- Adaptive Random Testing (ART)
 - Distance between inputs, and the 50% upper bound
- Fuzz Testing
 - Generation-based (grammars, Csmith) / mutation-based (bytes, ASTs)
 - **Coverage-guided** fuzzing, and directed fuzzing
 - Sanitizers as oracles, and crash triage

Last week a fuzzer kept a mutant when it **covered something new**. That “something” was one particular answer: a bitmap of **branches taken**.

- A **coverage criterion** needs a **structure** to be defined on – a **graph** or a **logic expression**.
- It turns that structure into a finite set of **test requirements**, so “how much have we tested?” becomes a **number**.
- And then the honest question: does a bigger number mean **better tests**? The same question comes back for neural networks.

1. Graph Coverage

- Structural Coverage

- Data-Flow Coverage

- Subsumption Relationships

2. Logic Coverage

- Simple Logic Expression Coverage

- Active Clause Coverage

- Inactive Clause Coverage

- Subsumption Relationships

3. Measuring Coverage

4. Is Coverage a Good Criterion?

5. Coverage Beyond Code

- Neuron Coverage

- Feature-Sensitive Coverage

1. Graph Coverage

- Structural Coverage

- Data-Flow Coverage

- Subsumption Relationships

2. Logic Coverage

- Simple Logic Expression Coverage

- Active Clause Coverage

- Inactive Clause Coverage

- Subsumption Relationships

3. Measuring Coverage


4. Is Coverage a Good Criterion?

5. Coverage Beyond Code

- Neuron Coverage

- Feature-Sensitive Coverage

```
int f(int a, int b) {  
    int x = 42;           // 0  
    if (a > 0) g();        // 1  
    else h();              // 2  
    int y;  
    if (b > 0)              // 3  
        y = x * 2;         // 4  
    else                    // 5  
        y = x - 5;         // 5  
    return y;              // 6  
}
```




- A **control flow graph (CFG)** is one of many: **call graphs**, **finite state machines**, class hierarchies, spec structures.
- Define coverage **once** on graphs, and it applies to all of them.

Definition (Graph)

A **graph** $G = (N, E, N_s, N_f)$ consists of a set of **nodes** N , a set of **edges** $E \subseteq N \times N$, a set of **start nodes** $N_s \subseteq N$, and a set of **final nodes** $N_f \subseteq N$.

- A **path** $p = (n_0, \dots, n_k) \in N^*$ satisfies $(n_i, n_{i+1}) \in E$ for $0 \leq i < k$; its **length** is $|p| = k$; and P_G is the set of all paths in G .
- A **subpath** $q \preceq p$ is a **contiguous** subsequence:
 $\exists 0 \leq i \leq j \leq k. q = (n_i, \dots, n_j)$.
- A **test path** starts at a start node and ends at a final node
 $(n_0 \in N_s \wedge n_k \in N_f)$ – it is what an **execution** of a test case looks like. We write $\text{path}(t)$ for the test path of a test case t , and $\text{path}(T)$ for those of a test suite T .

Consider a path $p = [0, 2, 4, 5, 4, 6]$ in the graph G .



- $|p| = 5$, and p is a **test path** ($0 \in N_s$, $6 \in N_f$)
- $[2, 4, 5] \preceq p$, but $[2, 4, 6] \not\preceq p$ – **contiguous**, not just in order
- p **visits** nodes 0, 2, 4, 5, 6 and edges $(0, 2)$, $(2, 4)$, $(4, 5)$, $(5, 4)$, $(4, 6)$

Definition (Graph Coverage Criterion)

A **graph coverage criterion** $C_G = (R_G, \sim)$ for a graph G is

- a set of **test requirements (TRs)** R_G , and
 - a **cover relation** $\sim \subseteq P_G \times R_G$ between paths and test requirements.
-
- A test case t **covers** a TR r if its test path does:
 $t \sim r \iff \text{path}(t) \sim r$.
 - A test suite T **satisfies** C_G if it covers every TR.

$$T \vdash C_G \iff \forall r \in R_G. \exists t \in T. t \sim r$$

- Everything below is a choice of R_G and $\sim$: **structural** criteria use only nodes, edges, and paths; **data-flow** criteria use a graph annotated with **variables**.


Definition (Node Coverage (NC))


$R_G = N$, and a path p **covers** a node n if p visits n .

Definition (Edge Coverage (EC))

$R_G = E$, and a path p **covers** an edge (n, m) if p visits (n, m) .

NC and EC differ only when there is **both** an edge and another subpath between a pair of nodes – an **if** without an **else** is the smallest case: $[0, 2]$ visits every node but misses $(0, 1)$.





- **Node Coverage (NC)**

$$R_G = \{0, 1, 2, 3, 4, 5, 6\}$$

$$\text{Test Paths} = \{[0, 1, 2, 3, 6], [0, 1, 2, 4, 5, 4, 6]\}$$

- **Edge Coverage (EC)**

$$R_G = \{(0, 1), (0, 2), (1, 2), (2, 3), (2, 4), (3, 6), (4, 5), (5, 4), (4, 6)\}$$

$$\text{Test Paths} = \{[0, 1, 2, 3, 6], [0, 2, 4, 5, 4, 6]\}$$

- The NC suite is **not** EC-adequate: it never takes (0, 2).


Definition (k -Limiting Path Coverage (k -PC))

- $R_G = \{p \in P_G \mid |p| \leq k\}$
- $p \sim q \iff q \preceq p$

- The criteria we just defined are the small cases of this one:

$k = 0$: paths of length 0 = nodes	node coverage (NC)
$k = 1$: + paths of length 1 = edges	edge coverage (EC)
$k = 2$: + pairs of adjacent edges	edge-pair coverage (EPC)
$k = \infty$: every path	complete path coverage (CPC)

- Strictly, 1-PC is NC **and** EC together. They coincide as long as every node touches an edge – an **isolated** node is a TR of 1-PC that EC alone would miss.
- So k is a single **knob** trading cost for strength.



- **11** requirements of length 2 (plus the nodes and edges):

$[0, 1, 2], [0, 2, 3], [0, 2, 4], [1, 2, 3], [1, 2, 4], [2, 3, 6],$
 $[2, 4, 5], [2, 4, 6], [4, 5, 4], [5, 4, 5], [5, 4, 6]$

- Test Paths

$= \{[0, 1, 2, 3, 6], [0, 1, 2, 4, 6], [0, 2, 3, 6], [0, 2, 4, 5, 4, 5, 4, 6]\}$

- Note $[5, 4, 5]$: covering it needs $4 \rightarrow 5 \rightarrow 4 \rightarrow 5$, so EPC forces the loop body **twice**.


- One loop makes P_G **infinite**, so CPC is infeasible. Which finitely many paths should stand in for the infinitely many?
- **Specified path coverage (SPC)** lets the tester list them – and makes the criterion depend on the tester's expertise.
- Four decades of answers:
 - 1970s: execute each cycle once ($[4, 5, 4]$ above)
 - 1980s: execute each loop exactly once
 - 1990s: execute each loop 0 times, once, more than once
 - 2000s: **prime paths** – get all three for free

Definition (Simple Path)

A **simple path** has no repeated node, except that the first and the last may coincide. (No internal loop; a whole loop is a simple path.)

Definition (Prime Path)


A **prime path** is a simple path that is not a subpath of any other simple path – a **maximal** simple path.



23 simple paths, of which **8** are prime:

[0, 1, 3, 0], [0, 2, 3, 0], [1, 3, 0, 1], [1, 3, 0, 2],
[2, 3, 0, 1], [2, 3, 0, 2], [3, 0, 1, 3], [3, 0, 2, 3]

The other 15 are subpaths of these: [0], [0, 1], [0, 1, 3], ...



38 simple paths, **9** prime:

[0, 1, 2, 3, 6]

[0, 1, 2, 4, 5]

[0, 1, 2, 4, 6]

[0, 2, 3, 6]

[0, 2, 4, 5]

[0, 2, 4, 6]

[4, 5, 4]

[5, 4, 5]

[5, 4, 6]

[0, 2, 4, 6] executes the loop **0** times

[4, 5, 4] executes the loop **once**

[5, 4, 5] executes the loop **more than once**

Prime path coverage needs all three, so it handles loops without a special rule.

Definition (Prime Path Coverage (PPC))

$R_G = \{p \in P_G \mid p \text{ is a prime path}\}$, and $p \sim q \iff q \preceq p$.

Not every graph is a CFG. In a **state machine**, testing often focuses on **loops**.¹


Definition (Round-Trip Path)

A **round-trip path** is a prime path that starts and ends at the **same** node.

- **Complete round trip coverage (CRTC)** requires **every** round-trip path; **simple round trip coverage (SRTC)** requires **one** per node that has any.
- Both are **weaker** than prime path coverage: they only cover the loops.

¹[Binder'99] R. V. Binder. "Testing Object-Oriented Systems: Models, Patterns, and Tools." [ISSRE'12] N. E. Holt, R. Torkar, L. Briand, K. Hansen. "State-Based Testing."


Prime paths have **no internal loops** – but a real execution may take one on the way. Must we reject it?



- p **tours** q if $q \preceq p$. $[0, 1, 2, 4, 5]$ tours $[1, 2, 4]$
- p tours q **with sidetrips** if every **edge** of q is in p in the same order.
 $[0, 1, 2, \textcolor{red}{3}, \textcolor{red}{2}, 4, 5]$
- p tours q **with detours** if every **node** of q is in p in the same order.
 $[0, 1, 2, \textcolor{red}{3}, 4, 5]$
- A detour may **skip the edge under test** – $(2, 4)$ is never taken above. That is why sidetrips are the usual compromise.

- A test requirement is **infeasible** if **no input** satisfies it: dead code, or a subpath that needs contradictory conditions (`if (x > 0)` inside `if (x < 0)`).
- Deciding feasibility is **undecidable** in general – so “100% coverage” is usually **unreachable**, and we cannot even tell how unreachable.
- The **stronger** the criterion, the more infeasible TRs: PPC on a loop whose bound is a constant, EPC on a branch that never repeats. Disallowing sidetrips makes it worse.
- **In practice**: satisfy what you can **without** sidetrips, then **allow** sidetrips for the rest, then inspect the remainder by hand.

Structural criteria ask **where control went**. Data-flow criteria ask whether a computed **value** was ever **used**.



$$\begin{aligned}
 \text{def}(0) &= \{x\} & \text{use}(4) &= \{x\} \\
 \text{def}(4) &= \{y\} & \text{use}(5) &= \{x\} \\
 \text{def}(5) &= \{y\} & \text{use}(6) &= \{y\}
 \end{aligned}$$

- A **def** of a variable is a location where it is **assigned**; a **use** is a location where its value is **read**.
- Node coverage reaches node 4. It never checks that the 42 written at node 0 is the value node 4 reads.

Definition (DU-Pair)

A **du-pair** for x is a pair of locations (l, l') such that x is defined at l and used at l' .

Definition (Def-Clear Path)

A path from l to l' is **def-clear** with respect to x if x is **not** redefined on any node or edge strictly between them.

Definition (DU-Path)

A **du-path** for x is a **simple** path from a def of x to a use of x that is **def-clear** with respect to x .


- $du(n, n', x)$ – du-paths from n to n' ; $du(n, x) = \bigcup_{n'} du(n, n', x)$
- A test path **du-tours** a du-path if it tours it **and** the subpath taken stays def-clear – sidetrips are allowed only if they do not redefine x .

One def, one use, and many paths between them. The three criteria differ exactly on **how many of those paths** you must take.

	test requirements R_G	p covers it by du-touring
ADC (<i>all-defs</i>)	(n, x)	some $q \in du(n, x)$
AUC (<i>all-uses</i>)	(n, n', x)	some $q \in du(n, n', x)$
ADUPC (<i>all-du-paths</i>)	every du-path q	q itself

only those with at least one du-path ($du(\cdot) \neq \emptyset$)

- ADC: every def reaches **some** use. AUC: every def reaches **every** use it can. ADUPC: and by **every route**.
- So ADUPC is the **strongest** of the three and ADC the **weakest**.
When a tool says “data-flow coverage” it almost always means **AUC** – the other two are too weak and too expensive respectively.



- **All-Defs Coverage (ADC)** – one du-path out of each of the three defs

$$R_G = \{(0, x), (4, y), (5, y)\} \quad \text{Test Paths} = \{[0, 1, 3, 4, 6], [0, 1, 3, 5, 6]\}$$

- **All-Uses Coverage (AUC)** – every (def, use) pair

$$R_G = \{(0, 4, x), (0, 5, x), (4, 6, y), (5, 6, y)\}$$
$$\text{Test Paths} = \{[0, 1, 3, 4, 6], [0, 1, 3, 5, 6]\}$$


- **All-DU-Paths Coverage (ADUPC)** – both routes through nodes 1 and 2

$$R_G = \{[0, 1, 3, 4], [0, 1, 3, 5], [0, 2, 3, 4], [0, 2, 3, 5], [4, 6], [5, 6]\}$$
$$\text{Test Paths} = \{[0, 1, 3, 4, 6], [0, 1, 3, 5, 6], [0, 2, 3, 4, 6], [0, 2, 3, 5, 6]\}$$

Definition (Subsumption)

C_G **subsumes** C'_G if **every** test suite satisfying C_G also satisfies C'_G .

- EC subsumes NC; EPC subsumes EC; PPC subsumes EC.
On the data-flow side, ADUPC subsumes AUC subsumes ADC.
- But PPC does **not** subsume EPC: if n has a self-loop, EPC requires $[n, n, m]$, which is not simple and so never a prime path.
- CRTC and SRTC subsume neither – they see only loops.
- Subsumption is about **guarantees**, not about **finding bugs**. A PPC-adequate suite is not automatically better at detecting faults than an EC-adequate one. We come back to this later.



1. Graph Coverage

Structural Coverage

Data-Flow Coverage

Subsumption Relationships

2. Logic Coverage

Simple Logic Expression Coverage

Active Clause Coverage

Inactive Clause Coverage

Subsumption Relationships

3. Measuring Coverage

4. Is Coverage a Good Criterion?

5. Coverage Beyond Code

Neuron Coverage

Feature-Sensitive Coverage

$$p = \underbrace{(a < b)}_{\text{relational}} \vee \underbrace{f(z)}_{\text{boolean function}} \wedge \underbrace{D}_{\text{variable}} \wedge \underbrace{(m \geq n \times o)}_{\text{relational}}$$

- A **predicate** is an expression evaluating to a **boolean**; a **clause** is a predicate with **no logical operator** in it. The predicate above has **four** clauses.
- Operators: $\neg$, $\wedge$, $\vee$, $\oplus$, $\Rightarrow$, $\Leftrightarrow$.
- Node and edge coverage see **one bit** per **if** – which way it went. Logic coverage looks **inside** the condition.

Notation: P is the set of predicates, C_p the clauses of predicate p .

Definition (Predicate Coverage (PC))

For each $p \in P$, the TRs are p evaluating to **true** and to **false**. (Also called **decision coverage**; “branch coverage” in tools.)

Definition (Clause Coverage (CC))

For each clause $c \in C_p$, the TRs are c evaluating to **true** and to **false**. (Also called **condition coverage**.)

PC asks **two** questions per predicate; CC asks **two per clause**. Both are cheap, and neither implies the other.

$$p = ((a < b) \vee D) \wedge (m \geq n \times o)$$

PC

a	b	D	m	n	o	p
5	10	T	1	1	1	T
10	5	F	1	1	1	F

CC

a	b	D	m	n	o	$(a < b)$	D	$(m \geq n \times o)$
5	10	F	1	1	1	T	F	T
10	5	T	1	2	2	F	T	F

- Two tests each. This CC suite **happens** to satisfy PC as well (p is **true** then **false**) – but that is luck, not a guarantee.

- **PC does not subsume CC**: one test pair can flip the predicate while leaving a clause always **false** – and with **short-circuit** evaluation, never even evaluating it.
- **CC does not subsume PC**: this suite covers both values of A and of B , yet $A \vee B$ is **true** in both rows.

A	B	$A \vee B$
T	F	T
F	T	T

- The simplest fix is to require **every combination**.

Definition (Combinatorial Coverage (CoC))

For each $p \in P$, the TRs are **all combinations** of truth values of the clauses in C_p .

$$p = ((a < b) \vee D) \wedge (m \geq n \times o)$$

$(a < b)$	D	$(m \geq n \times o)$	p
T	T	T	T
T	T	F	F
T	F	T	T
T	F	F	F
F	T	T	T
F	T	F	F
F	F	T	F
F	F	F	F

2^N tests for N clauses. We want CoC's **guarantee** at CC's **price**.

Definition (Determination)

A clause $c \in C_p$ – the **major clause** – **determines** p if the remaining **minor clauses** have values such that **changing the value of c changes the value of p** .

So a pair of tests that differs **only in** c and gives a **different** p is direct evidence that c matters.

A	B	$A \vee B$
T	T	T
T	F	T
F	T	T
F	F	F

A	B	$A \wedge B$
T	T	T
T	F	F
F	T	F
F	F	F

A	B	$A \oplus B$
T	T	F
T	F	T
F	T	T
F	F	F

A	B	$A \Leftrightarrow B$
T	T	T
T	F	F
F	T	F
F	F	T

- A determines $A \vee B$ only when $B = \text{false}$, and $A \wedge B$ only when $B = \text{true}$.
- A **always** determines $A \oplus B$ and $A \Leftrightarrow B$ – the minor clause cannot mask it.

Definition (Active Clause Coverage (ACC))

For each $p \in P$ and each major clause $c \in C_p$, the TRs are the minor values that make c **determine** p , paired with $c = \text{true}$ and $c = \text{false}$.

- For $p = A \vee B$:
 - 1 $A = \text{true}, B = \text{false}$ (A determines p)
 - 2 $A = \text{false}, B = \text{false}$ (A determines p)
 - 3 $A = \text{false}, B = \text{true}$ (B determines p)
 - 4 ~~$A = \text{false}, B = \text{false}$~~ (same row as 2)

$N + 1$ tests instead of 2^N : **three** rows, not four.

- ACC is better known as **modified condition/decision coverage (MC/DC)**. It is the criterion regulators require for the software that **flies aircraft**.²

²[SEJ'94] J. J. Chilenski, S. P. Miller. "Applicability of Modified Condition/Decision Coverage to Software Testing."

Ambiguity: must the **minor** clauses keep the **same values** when the major clause flips?

$$p = A \vee (B \wedge C), \quad \text{major clause } A$$

A	B	C	p
T	F	T	T
F	F	F	F

is $C = \text{false}$ allowed here?

This question confused testers – and certification authorities – for years.
Three answers, three criteria:³

- **General** active clause coverage (GACC)
- **Restricted** active clause coverage (RACC)
- **Correlated** active clause coverage (CACC)

³[ISSRE'03] P. Ammann, J. Offutt, H. Huang. "Coverage Criteria for Logical Expressions."

Definition (General Active Clause Coverage (GACC))

The minor clauses do **not need** to be the same when the major clause is **true** and when it is **false**.

- GACC does **not subsume predicate coverage**. For $A \Leftrightarrow B$, each clause always determines p , so these two rows satisfy GACC for both A and B – yet p is **true** in both.

A	B	$A \Leftrightarrow B$
T	T	T
F	F	T

- GACC was meant to be **stronger** than predicate coverage. It is not even as strong. So we need a different rule.

RACC

A	B	C	$A \wedge (B \vee C)$
<i>T</i>	T	T	T
<i>F</i>	T	T	F
<i>T</i>	T	F	T
<i>F</i>	T	F	F
<i>T</i>	F	T	T
<i>F</i>	F	T	F

CACC

A	B	C	$A \wedge (B \vee C)$
<i>T</i>	T	T	T
<i>T</i>	T	F	T
<i>T</i>	F	T	T
<i>F</i>	T	T	F
<i>F</i>	T	F	F
<i>F</i>	F	T	F

- **RACC** (*restricted*) pairs rows: B and C must be identical across the flip, so only **three** pairs are legal.
- **CACC** (*correlated*) only requires p to flip, so **any** top row pairs with **any** bottom row – **nine** choices.
- **RACC** is what the aviation industry adopted. **CACC** is the later reading: weaker, easier to satisfy, and still enough to subsume predicate coverage.

Definition (Inactive Clause Coverage (ICC))


For each $p \in P$ and each major clause $c \in C_p$, the TRs are the minor values that make c **not determine** p , paired with $c = \text{true}$ and $c = \text{false}$.

The **dual** of ACC. ACC picks tests where c **must** change the answer; ICC picks tests where c **must not**.

The **specification** says $p = A \vee B$, but code computes $A \oplus B$:

A	B	$A \vee B$	$A \oplus B$	
T	F	T	T	ACC
F	F	F	F	ACC
T	T	T	F	ICC
F	T	T	T	ICC

- Same ambiguity: **GICC** / **RICC**. No “correlated” variant – CACC asks the minor clauses to make p **flip** with c , and under ICC p never flips.



1. Graph Coverage

Structural Coverage

Data-Flow Coverage

Subsumption Relationships

2. Logic Coverage

Simple Logic Expression Coverage

Active Clause Coverage

Inactive Clause Coverage

Subsumption Relationships

3. Measuring Coverage

4. Is Coverage a Good Criterion?

5. Coverage Beyond Code

Neuron Coverage

Feature-Sensitive Coverage

$$\text{coverage}(T, C_G) = \frac{|\{r \in R_G \mid \exists t \in T. t \sim r\}|}{|R_G|}$$

- A criterion is **satisfied** or not; a **percentage** is what people actually report. Every criterion above gives one.
- **Infeasible** TRs stay in the denominator, so 100% is usually out of reach – and we do not know what the real maximum is. Teams either remove those TRs by hand or pick a lower target such as 85%.
- The denominator is **not stable**: adding a line changes $|R_G|$, so the percentage moves without anyone writing a test.
- Three questions remain: **where** do the numbers come from, **what** do real tools count, and **what does the number mean?**

To know that node 4 ran, the program must **say so**. Insert probes:

```
int f(int a, int b) {  
    cov[0] = 1; int x = 42;  
    if (a > 0) { cov[1] = 1; g(); }  
    else      { cov[2] = 1; h(); }  
    cov[3] = 1; int y;  
    if (b > 0) { cov[4] = 1; y = x * 2; }  
    else      { cov[5] = 1; y = x - 5; }  
    cov[6] = 1; return y;  
}
```

- Probes can go in at the **source**, **IR**, **bytecode**, or **binary**: later is truer to what runs, earlier is easier to map back to lines.
- You do not need a probe on every node. If node 6 **post-dominates** node 3, one probe tells you about both.
- Coverage asks **whether** a block ran, not how often – so a probe can be **removed** once it fires.⁴

⁴[ISSTA'02] M. M. Tikir, J. K. Hollingsworth. "Efficient Instrumentation for Code Coverage Testing."

- Everything shipped – **gcov**, **llvm-cov**, **JaCoCo**, **Istanbul** – measures **node or edge** coverage. Prime paths and all-du-paths live in textbooks and research tools, not in your CI.
- “Branch coverage” in a tool means **predicate coverage** on each decision – not edge coverage of the whole CFG, and not MC/DC.
- A fuzzer goes further: AFL **hashes** edges into 64K buckets, so two edges can share one. It wants a **fast** signal, not an exact one – the criterion follows from the **purpose**.

1. Graph Coverage

Structural Coverage

Data-Flow Coverage

Subsumption Relationships

2. Logic Coverage

Simple Logic Expression Coverage

Active Clause Coverage

Inactive Clause Coverage

Subsumption Relationships

3. Measuring Coverage

4. Is Coverage a Good Criterion?

5. Coverage Beyond Code

Neuron Coverage

Feature-Sensitive Coverage

```
TEST(FibTest, Smoke) {  
    fib(10);           // 100% line and branch coverage of fib  
}
```

- Coverage measures what was **executed**, never what was **checked**. The **oracle** is a separate problem, and a later topic.
- Even with an oracle: a fault must be **reached**, the state must be **infected**, and the corruption must **propagate** to the output. Coverage certifies only the **first**.
- **Necessary, not sufficient**. So how much does the necessary condition buy?

- Bigger suites cover more **and** find more. Control for suite size and the correlation with fault detection is **weak** – and the studies still argue about how weak.⁵

Low coverage is reliable bad news; high coverage is weak good news.

- **As a filter.** Google computes it for a billion lines a day and shows it on the **changed lines** of a code review – never as a quota.⁶
Goodhart's law: a measure that becomes a target stops being a good measure.
- **As a search gradient.** Last week's fuzzer, and techniques still to come. Nobody reports the percentage; it is only a signal.
- **Not as a quality measure.** For that, seed faults and count kills – **mutation score**, which we take up later.

⁵[ICSE'14] L. Inozemtseva, R. Holmes. "Coverage Is Not Strongly Correlated with Test Suite Effectiveness." [ASE'20] Y. T. Chen et al. "Revisiting the Relationship between Fault Detection, Test Adequacy Criteria, and Test Set Size."

⁶[FSE'19] M. Ivanković, G. Petrović, R. Just, G. Fraser. "Code Coverage at Google."

1. Graph Coverage

Structural Coverage

Data-Flow Coverage

Subsumption Relationships

2. Logic Coverage

Simple Logic Expression Coverage

Active Clause Coverage

Inactive Clause Coverage

Subsumption Relationships


3. Measuring Coverage

4. Is Coverage a Good Criterion?

5. Coverage Beyond Code

Neuron Coverage

Feature-Sensitive Coverage



- A DNN classifier is a function $f_\theta : X \rightarrow Y$ whose behavior lives in the **weights**, not in control flow. There is nothing to branch on – so cover the **neurons** instead. Write $f_\theta(n, x)$ for the output of neuron n on input x ; the neuron is **activated** when that exceeds a threshold t .

Definition (Neuron Coverage (NC))

$$NC(T, t) = \frac{|\{n \in N \mid \exists x \in T. f_\theta(n, x) > t\}|}{|N|}$$

Plain NC **saturates**, so it was refined by splitting each neuron's training range into k sections (**KMNC**), or looking outside that range (**NBC**), or taking the top k per layer (**TKNC**) – **the same k knob again.**⁷

⁷[SOSP'17] K. Pei, Y. Cao, J. Yang, S. Jana. “DeepXplore.” [ASE'18] L. Ma et al. “DeepGauge.”

Unfortunately, no!

- Generate inputs that **maximize** NC and see what you get: **fewer** defects detected, **less natural** inputs, and **more biased** predictions. Higher NC made the suite **worse**.⁸
- Independently, the criteria correlate poorly with model quality.⁹
- A better-founded alternative measures **surprise**: how far an input's activation trace is from the traces seen during training. It correlates with misclassification and gives a usable sampling target.¹⁰
- Same story as code coverage: **a criterion has to be validated, not just defined.**

⁸[FSE'20] F. Harel-Canada, L. Wang, M. A. Gulzar, Q. Gu, M. Kim. "Is Neuron Coverage a Meaningful Measure for Testing Deep Neural Networks?"

⁹[FSE'20] S. Yan, G. Tao, X. Liu, J. Zhai, S. Ma, L. Xu, X. Zhang. "Correlations between Deep Neural Network Model Coverage Criteria and Model Quality."


¹⁰[ICSE'19] J. Kim, R. Feldt, S. Yoo. "Guiding Deep Learning System Testing Using Surprise Adequacy."

```
a = [10, 20, 30];  
a.at(-1); // 30  
a.at(5); // undefined
```

23.1.3.1 Array.prototype.at (*index*)


1. Let *O* be ? ToObject(*this* value).
2. Let *len* be ? LengthOfArrayLike(*O*).
3. Let *relativeIndex* be ? ToIntegerOrInfinity(*index*).
4. If *relativeIndex* ≥ 0 , then
 - a. Let *k* be *relativeIndex*.
5. Else,
 - a. Let *k* be *len* + *relativeIndex*.
6. If *k* < 0 or *k* $\geq$ *len*, return **undefined**.
7. Return ? Get(*O*, ! ToString($\mathbb{F}(k)$)).

- ECMA-262 defines JavaScript as **pseudocode algorithms**. The two calls take **different paths** through this one: `at(-1)` runs $5 \rightarrow 7$, `at(5)` runs $4 \rightarrow 6$.




- **JEST**: 1,700 programs, **87.70%** of ES11 semantics – **44 engine bugs, 27 spec bugs**.¹¹

¹¹[\[ICSE'21\]](#) J. Park, S. An, D. Youn, G. Kim, S. Ryu. "JEST: N+1-Version Differential Testing of Both JavaScript Engines and Specification."



- Feature-Sensitive (FS)** coverage criterion **divides** the given TRs with the **innermost enclosing** language **features**



- 143** bugs in eight implementations.¹² Selective FS: **70.7%** fewer tests but preserves 53 of 57 bugs.¹³

¹²[PLDI'23] J. Park, D. Youn, K. Lee, S. Ryu. "Feature-Sensitive Coverage for Conformance Testing of Programming Language Implementations."

¹³[TOSEM'26] K. Lee, S. Kim, J. Park, S. Ryu. "Selective Feature-Sensitive Coverage for Conformance Testing of Programming Language Implementations."

1. Graph Coverage

- Structural Coverage

- Data-Flow Coverage

- Subsumption Relationships

2. Logic Coverage

- Simple Logic Expression Coverage

- Active Clause Coverage

- Inactive Clause Coverage

- Subsumption Relationships

3. Measuring Coverage

4. Is Coverage a Good Criterion?

5. Coverage Beyond Code

- Neuron Coverage

- Feature-Sensitive Coverage

- Build a **coverage profiler** that measures which parts of a program are executed by a test suite.
- Please see this document on GitHub:

<https://github.com/ku-plrg-classroom/js-cov>

- The due date is **23:59 on September 21 (Mon.)**, and please submit it to [LMS](#).

- Search Based Software Testing

Jihyeok Park

jihyeok_park@korea.ac.kr

<https://plrg.korea.ac.kr>