

Lecture 6 – Regression Testing

AAA705: Software Testing and Quality Assurance

Jihyeok Park



2026 Fall

- Symbolic Execution
 - Symbols in place of inputs; a path condition per path; an SMT solver returns the input
 - $hash(\alpha) = \beta$: no theory, so the solver says **UNKNOWN**
- Dynamic Symbolic Execution (DSE)
 - Run on concrete values too, and use them where the solver cannot
 - Search heuristics – DFS, CFDS, CGS, learned; loops and data structures
- Hybrid Analysis – blended, heap snapshots, Concerto, dynamic shortcuts

- **So far:** how to **make** a test – random, coverage-guided, search-based, symbolic. Every test we made stays in the suite.

- **So far**: how to **make** a test – random, coverage-guided, search-based, symbolic. Every test we made stays in the suite.
- The program **changes** every day, and the suite only **grows**. Running all of it after every change becomes too slow.

- **So far**: how to **make** a test – random, coverage-guided, search-based, symbolic. Every test we made stays in the suite.
- The program **changes** every day, and the suite only **grows**. Running all of it after every change becomes too slow.
- **Today**: which tests to **drop**, which to **rerun** after a change, and in which **order**.

1. Regression Testing

- Regression Fault

- Test Suite Minimization

- Test Case Selection

- Test Case Prioritization

- Regression Testing in Practice


iPhone, iOS 17.4.1 (April 2024)¹

If your iPhone alarm stopped working in iOS 17.4.1, here's a fix

How-To

Apr 30, 2024

Apple is working to patch the bug but a simple tweak in Settings should help.



¹[Macworld, Apr. 30, 2024] "If your iPhone alarm stopped working in iOS 17.4.1, here's a fix." Apple: "aware of an issue causing some alarms not play a sound and . . . working on a fix."

Windows 11, June 2026 update (KB5094126)²

This article was published on June 21, 2026

DATA AND SECURITY

Microsoft's June update fixed 208 security flaws and introduced a cascade of new bugs across all Windows versions

KB5094126 patched 208 security vulnerabilities on June 9, but the update has triggered Recycle Bin display glitches, BitLocker lockouts on enterprise hardware, OneDrive failures, and system freezes, with the first fix not expected until July 14

June 21, 2026 - 1:59 pm

²[[TNW, June 21, 2026](#)] "Microsoft's June update fixed 208 security flaws and introduced a cascade of new bugs across all Windows versions."

- **Regression fault** is a fault that occurs when a **change** in the software introduces a new defect or reactivates a defect that had been previously fixed.
- The term **regression** refers to the fact that the software has **regressed** (gone backwards) to an earlier bad state.

- You are releasing a **new version** of your software.

- You are releasing a **new version** of your software.
- You have tried to thoroughly **test the new features**.

- You are releasing a **new version** of your software.
- You have tried to thoroughly **test the new features**.
- You want to check if you have created any **regression faults**.

- You are releasing a **new version** of your software.
- You have tried to thoroughly **test the new features**.
- You want to check if you have created any **regression faults**.
- How would you do that?

- The simplest way is to **retest all** the test cases.

- The simplest way is to **retest all** the test cases.
- You need to run all tests not only for the **new features** but also for the **old existing features**.

- The simplest way is to **retest all** the test cases.
- You need to run all tests not only for the **new features** but also for the **old existing features**.
- Its main disadvantage is that it is **time-consuming** to run all the test cases for every new version.

- The simplest way is to **retest all** the test cases.
- You need to run all tests not only for the **new features** but also for the **old existing features**.
- Its main disadvantage is that it is **time-consuming** to run all the test cases for every new version.
- It is **critical** in the modern software development process because software is **continuously and rapidly changing**.

*“For example, one of our industrial collaborators reports that for one of its products of about **20,000 lines of code**, the entire test suite requires **seven weeks** to run.”*

– Rothermel et al.³

³[TSE'01] G. Rothermel, R. H. Untch, C. Chu, M. J. Harrold. “Prioritizing Test Cases for Regression Testing.”

*“For example, one of our industrial collaborators reports that for one of its products of about **20,000 lines of code**, the entire test suite requires **seven weeks** to run.”*

– Rothermel et al.³

The **test suite** becomes **larger** and **larger** as the software evolves with the following factors:

- Long Product History
- Different Configurations
- Types of Test Cases

³[TSE'01] G. Rothermel, R. H. Untch, C. Chu, M. J. Harrold. “Prioritizing Test Cases for Regression Testing.”

Many techniques have been developed in order to cope with the high cost of retesting all test cases. They can be categorized into:

- **Test Suite Minimization**
- **Test Case Selection**
- **Test Case Prioritization**

- **Problem** – Regression test suite is **too large**

- **Problem** – Regression test suite is **too large**
- **Idea** – There might be some test cases that are **redundant**

- **Problem** – Regression test suite is **too large**
- **Idea** – There might be some test cases that are **redundant**
- **Solution** – **Minimize** regression test suite by removing all the redundant test cases


- **Problem** – Regression test suite is **too large**
- **Idea** – There might be some test cases that are **redundant**
- **Solution** – **Minimize** regression test suite by removing all the redundant test cases
- Then, what is the definition of a **redundant** test case?

One possible way is to use the **coverage** information.


(e.g., Statements, MC/DC, All-DU-Paths, etc.)

TC	r_1	r_2	r_3	$\dots$
t_1	✓	✓		$\dots$
t_2		✓		$\dots$
t_3		✓	✓	$\dots$
t_4			✓	$\dots$
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\ddots$


How to **minimize** the test suite as much as possible while **preserving** the coverage information (the set of covered test requirements)?



- We can model the **test suite minimization** problem as a **set cover** problem.



- We can model the **test suite minimization** problem as a **set cover** problem.
- Unfortunately, the set cover problem is **NP-complete**, meaning that there is no known polynomial-time algorithm to solve it.



- We can model the **test suite minimization** problem as a **set cover** problem.
- Unfortunately, the set cover problem is **NP-complete**, meaning that there is no known polynomial-time algorithm to solve it.
- However, there are many **heuristic** algorithms that can be used to solve the test suite minimization problem.

Algorithm Greedy Minimization

```
1: function GREEDYMINIMIZATION( $T, R$ )
2:    $T' \leftarrow \emptyset$ 
3:   while true do
4:      $t \leftarrow \operatorname{argmax}_{t \in T} |R \cap \operatorname{TR}(t)|$ 
5:     if  $|R \cap \operatorname{TR}(t)| = 0$  then
6:       break
7:      $T' \leftarrow T' \cup \{t\}$ 
8:      $R \leftarrow R \setminus \operatorname{TR}(t)$ 
9:   return  $T'$ 
```

Algorithm Greedy Minimization

```

1: function GREEDYMINIMIZATION( $T, R$ )
2:    $T' \leftarrow \emptyset$ 
3:   while true do
4:      $t \leftarrow \operatorname{argmax}_{t \in T} |R \cap \operatorname{TR}(t)|$ 
5:     if  $|R \cap \operatorname{TR}(t)| = 0$  then
6:       break
7:      $T' \leftarrow T' \cup \{t\}$ 
8:      $R \leftarrow R \setminus \operatorname{TR}(t)$ 
9:   return  $T'$ 
    
```

TC	r_1	r_2	r_3	r_4
t_1	✓	✓		
t_2		✓		✓
t_3		✓	✓	✓
t_4			✓	
t_5			✓	✓

$$T' = \emptyset$$

$$R = \{r_1, r_2, r_3, r_4\}$$

Algorithm Greedy Minimization

```

1: function GREEDYMINIMIZATION( $T, R$ )
2:    $T' \leftarrow \emptyset$ 
3:   while true do
4:      $t \leftarrow \operatorname{argmax}_{t \in T} |R \cap \text{TR}(t)|$ 
5:     if  $|R \cap \text{TR}(t)| = 0$  then
6:       break
7:      $T' \leftarrow T' \cup \{t\}$ 
8:      $R \leftarrow R \setminus \text{TR}(t)$ 
9:   return  $T'$ 
    
```

TC	r_1	r_2	r_3	r_4
t_1	✓	✓		
t_2		✓		✓
t_3		✓	✓	✓
t_4			✓	
t_5			✓	✓

$$T' = \emptyset$$

$$R = \{r_1, r_2, r_3, r_4\}$$

$$T' = \{t_3\}$$

$$R = \{r_1\}$$

Algorithm Greedy Minimization

```

1: function GREEDYMINIMIZATION( $T, R$ )
2:    $T' \leftarrow \emptyset$ 
3:   while true do
4:      $t \leftarrow \operatorname{argmax}_{t \in T} |R \cap \text{TR}(t)|$ 
5:     if  $|R \cap \text{TR}(t)| = 0$  then
6:       break
7:      $T' \leftarrow T' \cup \{t\}$ 
8:      $R \leftarrow R \setminus \text{TR}(t)$ 
9:   return  $T'$ 
    
```

TC	r_1	r_2	r_3	r_4
t_1	✓	✓		
t_2		✓		✓
t_3		✓	✓	✓
t_4			✓	
t_5			✓	✓

$$T' = \emptyset$$

$$R = \{r_1, r_2, r_3, r_4\}$$

$$T' = \{t_3\}$$

$$R = \{r_1\}$$

$$T' = \{t_3, t_1\}$$

$$R = \emptyset$$

- However, in fact, test cases have different **costs** to run.

- However, in fact, test cases have different **costs** to run.
- Is $\{t_3, t_1\}$ still the best solution?

TC	r_1	r_2	r_3	r_4	<i>Time</i>
t_1	✓	✓			3
t_2		✓		✓	5
t_3		✓	✓	✓	10
t_4			✓		2
t_5			✓	✓	8

Algorithm Greedy Minimization with Cost

```
1: function GREEDYMINIMIZATION( $T, R$ )
2:    $T' \leftarrow \emptyset$ 
3:   while true do
4:      $t \leftarrow \operatorname{argmax}_{t \in T} \frac{|R \cap \text{TR}(t)|}{\text{Time}(t)}$ 
5:     if  $|R \cap \text{TR}(t)| = 0$  then
6:       break
7:      $T' \leftarrow T' \cup \{t\}$ 
8:      $R \leftarrow R \setminus \text{TR}(t)$ 
9:   return  $T'$ 
```

Algorithm Greedy Minimization with Cost

```

1: function GREEDYMINIMIZATION( $T, R$ )
2:    $T' \leftarrow \emptyset$ 
3:   while true do
4:      $t \leftarrow \operatorname{argmax}_{t \in T} \frac{|R \cap \text{TR}(t)|}{\text{Time}(t)}$ 
5:     if  $|R \cap \text{TR}(t)| = 0$  then
6:       break
7:      $T' \leftarrow T' \cup \{t\}$ 
8:      $R \leftarrow R \setminus \text{TR}(t)$ 
9:   return  $T'$ 
  
```

TC	r_1	r_2	r_3	r_4	Time
t_1	✓	✓			3
t_2		✓		✓	5
t_3		✓	✓	✓	10
t_4			✓		2
t_5			✓	✓	8

 $T' = \emptyset$
 $R = \{r_1, r_2, r_3, r_4\}$

Algorithm Greedy Minimization with Cost

```

1: function GREEDYMINIMIZATION( $T, R$ )
2:    $T' \leftarrow \emptyset$ 
3:   while true do
4:      $t \leftarrow \operatorname{argmax}_{t \in T} \frac{|R \cap \text{TR}(t)|}{\text{Time}(t)}$ 
5:     if  $|R \cap \text{TR}(t)| = 0$  then
6:       break
7:      $T' \leftarrow T' \cup \{t\}$ 
8:      $R \leftarrow R \setminus \text{TR}(t)$ 
9:   return  $T'$ 

```

TC	r_1	r_2	r_3	r_4	Time
t_1	✓	✓			3
t_2		✓		✓	5
t_3		✓	✓	✓	10
t_4			✓		2
t_5			✓	✓	8

$$T' = \emptyset$$

$$R = \{r_1, r_2, r_3, r_4\}$$

$$T' = \{t_1\}$$

$$R = \{r_3, r_4\}$$

Algorithm Greedy Minimization with Cost

```

1: function GREEDYMINIMIZATION( $T, R$ )
2:    $T' \leftarrow \emptyset$ 
3:   while true do
4:      $t \leftarrow \operatorname{argmax}_{t \in T} \frac{|R \cap \text{TR}(t)|}{\text{Time}(t)}$ 
5:     if  $|R \cap \text{TR}(t)| = 0$  then
6:       break
7:      $T' \leftarrow T' \cup \{t\}$ 
8:      $R \leftarrow R \setminus \text{TR}(t)$ 
9:   return  $T'$ 
  
```

TC	r_1	r_2	r_3	r_4	Time
t_1	✓	✓			3
t_2		✓		✓	5
t_3		✓	✓	✓	10
t_4			✓		2
t_5			✓	✓	8

$$T' = \emptyset$$

$$R = \{r_1, r_2, r_3, r_4\}$$

$$T' = \{t_1\}$$

$$R = \{r_3, r_4\}$$

$$T' = \{t_1, t_4\}$$

$$R = \{r_4\}$$

Algorithm Greedy Minimization with Cost

```

1: function GREEDYMINIMIZATION( $T, R$ )
2:    $T' \leftarrow \emptyset$ 
3:   while true do
4:      $t \leftarrow \operatorname{argmax}_{t \in T} \frac{|R \cap \text{TR}(t)|}{\text{Time}(t)}$ 
5:     if  $|R \cap \text{TR}(t)| = 0$  then
6:       break
7:      $T' \leftarrow T' \cup \{t\}$ 
8:      $R \leftarrow R \setminus \text{TR}(t)$ 
9:   return  $T'$ 
    
```

TC	r_1	r_2	r_3	r_4	Time
t_1	✓	✓			3
t_2		✓		✓	5
t_3		✓	✓	✓	10
t_4			✓		2
t_5			✓	✓	8

$$T' = \emptyset$$

$$R = \{r_1, r_2, r_3, r_4\}$$

$$T' = \{t_1\}$$

$$R = \{r_3, r_4\}$$

$$T' = \{t_1, t_4\}$$

$$R = \{r_4\}$$

$$T' = \{t_1, t_4, t_2\} \quad R = \emptyset$$

- Another possible approach is to use the **score** information and keep the **best test case** for each test requirement according to the score.

- Another possible approach is to use the **score** information and keep the **best test case** for each test requirement according to the score.
- The **score** of each test case can be defined in various ways.

- Another possible approach is to use the **score** information and keep the **best test case** for each test requirement according to the score.
- The **score** of each test case can be defined in various ways.
- For example, consider a **conformance test suite** for JavaScript engines: each test case is a JavaScript program with assertions.

- Another possible approach is to use the **score** information and keep the **best test case** for each test requirement according to the score.
- The **score** of each test case can be defined in various ways.
- For example, consider a **conformance test suite** for JavaScript engines: each test case is a JavaScript program with assertions.
- Then, which test case is better to keep?

- Another possible approach is to use the **score** information and keep the **best test case** for each test requirement according to the score.
- The **score** of each test case can be defined in various ways.
- For example, consider a **conformance test suite** for JavaScript engines: each test case is a JavaScript program with assertions.
- Then, which test case is better to keep?
- We can define the score of each test case based on the **complexity** of the test case (e.g., **size of the program**, etc.) because a simpler test case is more understandable and maintainable.

Algorithm Greedy Minimization with Score

```
1: function GREEDYMINIMIZATION( $T, R$ )
2:    $M \leftarrow \emptyset$ 
3:   for  $t \in T$  do
4:     for  $r \in \text{TR}(t)$  do
5:       if  $r \notin M \vee \text{Score}(t) > \text{Score}(M[r])$  then
6:          $M[r] \leftarrow t$ 
7:    $T' \leftarrow \{t \mid (r, t) \in M\}$ 
8:   return  $T'$ 
```

Algorithm Greedy Minimization with Score

```

1: function GREEDYMINIMIZATION( $T, R$ )
2:    $M \leftarrow \emptyset$ 
3:   for  $t \in T$  do
4:     for  $r \in \text{TR}(t)$  do
5:       if  $r \notin M \vee \text{Score}(t) > \text{Score}(M[r])$  then
6:          $M[r] \leftarrow t$ 
7:    $T' \leftarrow \{t \mid (r, t) \in M\}$ 
8:   return  $T'$ 
    
```

TC	r_1	r_2	r_3	r_4	Score
t_1	✓	✓			4
t_2		✓		✓	3
t_3		✓	✓	✓	2
t_4			✓		9
t_5			✓	✓	5

$$M = \left\{ \begin{array}{l} \\ \\ \\ \\ \end{array} \right.$$

Algorithm Greedy Minimization with Score

```

1: function GREEDYMINIMIZATION( $T, R$ )
2:    $M \leftarrow \emptyset$ 
3:   for  $t \in T$  do
4:     for  $r \in \text{TR}(t)$  do
5:       if  $r \notin M \vee \text{Score}(t) > \text{Score}(M[r])$  then
6:          $M[r] \leftarrow t$ 
7:    $T' \leftarrow \{t \mid (r, t) \in M\}$ 
8:   return  $T'$ 
  
```

TC	r_1	r_2	r_3	r_4	Score
t_1	✓	✓			4
t_2		✓		✓	3
t_3		✓	✓	✓	2
t_4			✓		9
t_5			✓	✓	5

$$M = \left\{ \begin{array}{l} r_1 \mapsto t_1 \end{array} \right.$$

Algorithm Greedy Minimization with Score

```
1: function GREEDYMINIMIZATION( $T, R$ )
2:    $M \leftarrow \emptyset$ 
3:   for  $t \in T$  do
4:     for  $r \in \text{TR}(t)$  do
5:       if  $r \notin M \vee \text{Score}(t) > \text{Score}(M[r])$  then
6:          $M[r] \leftarrow t$ 
7:    $T' \leftarrow \{t \mid (r, t) \in M\}$ 
8:   return  $T'$ 
```

TC	r_1	r_2	r_3	r_4	Score
t_1	✓	✓			4
t_2		✓		✓	3
t_3		✓	✓	✓	2
t_4			✓		9
t_5			✓	✓	5

$$M = \begin{cases} r_1 \mapsto t_1 \\ r_2 \mapsto t_1 \end{cases}$$

Algorithm Greedy Minimization with Score

```

1: function GREEDYMINIMIZATION( $T, R$ )
2:    $M \leftarrow \emptyset$ 
3:   for  $t \in T$  do
4:     for  $r \in \text{TR}(t)$  do
5:       if  $r \notin M \vee \text{Score}(t) > \text{Score}(M[r])$  then
6:          $M[r] \leftarrow t$ 
7:    $T' \leftarrow \{t \mid (r, t) \in M\}$ 
8:   return  $T'$ 
  
```

TC	r_1	r_2	r_3	r_4	Score
t_1	✓	✓			4
t_2		✓		✓	3
t_3		✓	✓	✓	2
t_4			✓		9
t_5			✓	✓	5

$$M = \begin{cases} r_1 \mapsto t_1 \\ r_2 \mapsto t_1 \\ r_3 \mapsto t_4 \end{cases}$$

Algorithm Greedy Minimization with Score

```

1: function GREEDYMINIMIZATION( $T, R$ )
2:    $M \leftarrow \emptyset$ 
3:   for  $t \in T$  do
4:     for  $r \in \text{TR}(t)$  do
5:       if  $r \notin M \vee \text{Score}(t) > \text{Score}(M[r])$  then
6:          $M[r] \leftarrow t$ 
7:    $T' \leftarrow \{t \mid (r, t) \in M\}$ 
8:   return  $T'$ 
  
```

TC	r_1	r_2	r_3	r_4	Score
t_1	✓	✓			4
t_2		✓		✓	3
t_3		✓	✓	✓	2
t_4			✓		9
t_5			✓	✓	5

$$M = \left\{ \begin{array}{l} r_1 \mapsto t_1 \\ r_2 \mapsto t_1 \\ r_3 \mapsto t_4 \\ r_4 \mapsto t_5 \end{array} \right\}$$

Algorithm Greedy Minimization with Score

```

1: function GREEDYMINIMIZATION( $T, R$ )
2:    $M \leftarrow \emptyset$ 
3:   for  $t \in T$  do
4:     for  $r \in \text{TR}(t)$  do
5:       if  $r \notin M \vee \text{Score}(t) > \text{Score}(M[r])$  then
6:          $M[r] \leftarrow t$ 
7:    $T' \leftarrow \{t \mid (r, t) \in M\}$ 
8:   return  $T'$ 
  
```

TC	r_1	r_2	r_3	r_4	Score
t_1	✓	✓			4
t_2		✓		✓	3
t_3		✓	✓	✓	2
t_4			✓		9
t_5			✓	✓	5

$$M = \left\{ \begin{array}{l} r_1 \mapsto t_1 \\ r_2 \mapsto t_1 \\ r_3 \mapsto t_4 \\ r_4 \mapsto t_5 \end{array} \right\}$$

$$T' = \{t_1, t_4, t_5\}$$

We can minimize the test suite even during the **test case generation**.⁴

⁴[\[ICSE'21\]](#) J. Park, S. An, D. Youn, G. Kim, S. Ryu. "*JEST: N+1-version Differential Testing of Both JavaScript Engines and Specification.*"

We can minimize the test suite even during the **test case generation**.⁴

RelationalExpression : *RelationalExpression* **<=** *ShiftExpression*

1. Let *lref* be ? **Evaluation** of *RelationalExpression*.
2. Let *lval* be ? **GetValue**(*lref*).
3. Let *rref* be ? **Evaluation** of *ShiftExpression*.
4. Let *rval* be ? **GetValue**(*rref*).
5. Let *r* be ? **IsLessThan**(*rval*, *lval*, **false**).
6. If *r* is either **true** or **undefined**, return **false**. Otherwise, return **true**.

⁴**[ICSE'21]** J. Park, S. An, D. Youn, G. Kim, S. Ryu. "JEST: N+1-version Differential Testing of Both JavaScript Engines and Specification."

Tests generated from the spec: keep, for each spec feature, the **smallest** program that covers it.⁵

RelationalExpression : *RelationalExpression* <= *ShiftExpression*

1. Let *lref* be ? **Evaluation** of *RelationalExpression*.
2. Let *lval* be ? **GetValue**(*lref*).
3. Let *rref* be ? **Evaluation** of *ShiftExpression*.
4. Let *rval* be ? **GetValue**(*rref*). ~~-----~~ **1 <= Symbol()**
5. Let *r* be **?IsLessThan**(*rval*, *lval*, **false**).
6. If *r* is either **true** or **undefined**, return **false**. Otherwise, return **true**.

⁵[ICSE'21] J. Park, S. An, D. Youn, G. Kim, S. Ryu. "JEST: N+1-version Differential Testing of Both JavaScript Engines and Specification."

What if the test suite minimization problem has **multiple objectives**?

What if the test suite minimization problem has **multiple objectives**?

- “I have **only 5 hours** as a deadline to run the test suite.”

What if the test suite minimization problem has **multiple objectives**?

- “I have **only 5 hours** as a deadline to run the test suite.”
- “I need to cover **more than one coverage** criterion.”

What if the test suite minimization problem has **multiple objectives**?

- “I have **only 5 hours** as a deadline to run the test suite.”
- “I need to cover **more than one coverage** criterion.”
- “I want to limit the **fault detection loss** of the minimized suite.”

$$\left(1 - \frac{\# \text{ Faults Detected by the Minimized Test Suite}}{\# \text{ Faults Detected by the Original Test Suite}}\right) \times 100\%$$

- The simplest way: **add the objectives up** with weights, and go back to a single-objective problem.

$$o' = \alpha_1 \times o_1 + \alpha_2 \times o_2 + \cdots + \alpha_n \times o_n \quad \sum_i \alpha_i = 1$$

- The simplest way: **add the objectives up** with weights, and go back to a single-objective problem.

$$o' = \alpha_1 \times o_1 + \alpha_2 \times o_2 + \cdots + \alpha_n \times o_n \quad \sum_i \alpha_i = 1$$

- Then, the greedy algorithm picks **one** test suite – the best one **for these weights**.

- The simplest way: **add the objectives up** with weights, and go back to a single-objective problem.

$$o' = \alpha_1 \times o_1 + \alpha_2 \times o_2 + \cdots + \alpha_n \times o_n \quad \sum_i \alpha_i = 1$$


- Then, the greedy algorithm picks **one** test suite – the best one **for these weights**.
- But who decides α ? Coverage 0.7 and time 0.3? Or 0.5 and 0.5? Each choice gives a **different** suite, and there is no way to tell which weights are right.

- A suite is **dominated** if another suite is at least as good in **every** objective and better in **one**.

- A suite is **dominated** if another suite is at least as good in **every** objective and better in **one**.
- The suites that **nobody dominates** are **Pareto efficient**. Together they form the **Pareto front**.

- A suite is **dominated** if another suite is at least as good in **every** objective and better in **one**.
- The suites that **nobody dominates** are **Pareto efficient**. Together they form the **Pareto front**.
- Instead of one suite, we get the **whole front** – and decide **later**.

- A suite is **dominated** if another suite is at least as good in **every** objective and better in **one**.
- The suites that **nobody dominates** are **Pareto efficient**. Together they form the **Pareto front**.
- Instead of one suite, we get the **whole front** – and decide **later**.
- “I have **5 hours**”: take the front, and pick the point **under 5 hours** with the highest coverage.



- One yes/no per test ($t_i = 1$ keep it, 0 drop it), and we keep as few as possible:

$$\min t_1 + t_2 + \cdots + t_n$$

- What we may not break – one line per test requirement r :

$$t_2, t_5, \text{ or } t_7 \text{ covers } r \quad \implies \quad t_2 + t_5 + t_7 \geq 1$$

- A sum to make small plus “at least one” lines: that is an **integer linear program (ILP)**. It is our set cover, written down – greedy was its quick approximation. Faults? One more line of the same shape.

⁶[TOSEM'25] S. Gu, A. Mesbah. “Scalable Similarity-Aware Test Suite Minimization with Reinforcement Learning.”

- One yes/no per test ($t_i = 1$ keep it, 0 drop it), and we keep as few as possible:

$$\min t_1 + t_2 + \cdots + t_n$$

- What we may not break – one line per test requirement r :

$$t_2, t_5, \text{ or } t_7 \text{ covers } r \quad \implies \quad t_2 + t_5 + t_7 \geq 1$$

- A sum to make small plus “at least one” lines: that is an **integer linear program (ILP)**. It is our set cover, written down – greedy was its quick approximation. Faults? One more line of the same shape.
- **TripRL** – the same ILP plus a penalty for each **similar pair**. Solvers stall past a few hundred tests, so a **reinforcement learning** agent picks them one at a time.⁶

⁶[TOSEM'25] S. Gu, A. Mesbah. “Scalable Similarity-Aware Test Suite Minimization with Reinforcement Learning.”

- **Problem** – Regression test suite is **too large**

- **Problem** – Regression test suite is **too large**
- **Idea** – Not all tests are **related** to the recent changes in the software.

- **Problem** – Regression test suite is **too large**
- **Idea** – Not all tests are **related** to the recent changes in the software.
- **Solution** – Precisely **select** the test cases that actually **execute** the changed parts of the software.

- **Problem** – Regression test suite is **too large**
- **Idea** – Not all tests are **related** to the recent changes in the software.
- **Solution** – Precisely **select** the test cases that actually **execute** the changed parts of the software.
- Then, how can we know which test cases are **related** to the recent changes?

- We have the original program P and the updated program P' .

- We have the original program P and the updated program P' .
- We need to keep track of the **execution trace** of the test cases in the original program P to know which parts of the program are executed by each test case.

- We have the original program P and the updated program P' .
- We need to keep track of the **execution trace** of the test cases in the original program P to know which parts of the program are executed by each test case.
- Then, we should know **which parts** of the program are **changed** in the updated program P' compared to the original program P .

- Most modern software is developed with a **version control system** (e.g., Git).

- Most modern software is developed with a **version control system** (e.g., Git).
- The easiest way is to use the **diff** command provided by the version control system to know the **changed parts** of the program.

<pre>191 def codeUnit: AbsCodeUnit 192 def const: AbsConst 193 def math: AbsMath 194 def simpleValue: AbsSimpleValue 195 def number: AbsNumber 196 def bigInt: AbsBigInt</pre>	<pre>194 def codeUnit: AbsCodeUnit 195 def const: AbsConst 196 def math: AbsMath 197 + def infinity: AbsInfinity 198 def simpleValue: AbsSimpleValue 199 def number: AbsNumber 200 def bigInt: AbsBigInt</pre>
---	---

▼ ↕ 203 ■■■ src/main/scala/esmeta/analyzer/domain/value/ValueTypeDomain.scala

☐ Viewed  ...

<pre>@@ -4,7 +4,8 @@ import esmeta.analyzer.* 4 import esmeta.analyzer.domain.* 5 import esmeta.cfg.Func 6 import esmeta.es.* 7 - import esmeta.ir.{COp, Name, VOp, 8 MOp}</pre>	<pre>4 import esmeta.analyzer.domain.* 5 import esmeta.cfg.Func 6 import esmeta.es.* 7 + import esmeta.ir.{COp, Name, VOp, 8 MOp, UOp} 8 + import esmeta.interpreter.Interpreter 9 import esmeta.parser.ESValueParser 10 import esmeta.state.* 11 import esmeta.spec.{Grammar => _, *}</pre>
--	---

A textual diff tells which **lines** changed, not which **tests** reach them.

A textual diff tells which **lines** changed, not which **tests** reach them.

P

```
function foo(x) {  
  if (x >= 0) {  
    return x;  
  } else {  
    return -x;  
  }  
}
```

P'

```
function foo(x) {  
  if (x >= 0) {  
    return x;  
  } else {  
    return 1 - x;  
  }  
}
```

Test Suite

- $t_1 : x = 1$
- $t_2 : x = -1$

A textual diff tells which **lines** changed, not which **tests** reach them.

P

```
function foo(x) {  
  if (x >= 0) {  
    return x;  
  } else {  
    return -x;  
  }  
}
```

P'

```
function foo(x) {  
  if (x >= 0) {  
    return x;  
  } else {  
    return 1 - x;  
  }  
}
```

Test Suite

- $t_1 : x = 1$
- $t_2 : x = -1$

Which test case should be **selected**?

A textual diff tells which **lines** changed, not which **tests** reach them.

P

```
function foo(x) {  
  if (x >= 0) {  
    return x;  
  } else {  
    return -x;  
  }  
}
```

P'

```
function foo(x) {  
  if (x >= 0) {  
    return x;  
  } else {  
    return 1 - x;  
  }  
}
```

Test Suite

- $t_1 : x = 1$
- $t_2 : x = -1$

Which test case should be **selected**?

- A file-level diff says foo changed: **both**.
- The trace of t_1 in P never reaches the changed statement: **only** t_2 .

- Selection techniques do **not consider** the **cost** of tests. Why?

- Selection techniques do **not consider** the **cost** of tests. Why?
- Because they focus on **safety**: no test related to the change may be missed, or a regression fault slips through.

- Selection techniques do **not consider** the **cost** of tests. Why?
- Because they focus on **safety**: no test related to the change may be missed, or a regression fault slips through.
- Realistically, we only consider whether each part of the program is **executed** by the test cases or not to check the safety according to their **execution traces**.

- **Data Collection** – Collecting the **execution traces** of the test cases is not easy. Especially, when the software is **large** and **complex**, and some programs are written in multiple languages.

- **Data Collection** – Collecting the **execution traces** of the test cases is not easy. Especially, when the software is **large** and **complex**, and some programs are written in multiple languages.
- **Non-executable modifications** – Some modifications are **not executable** (e.g., configuration changes, etc.)

- **Data Collection** – Collecting the **execution traces** of the test cases is not easy. Especially, when the software is **large** and **complex**, and some programs are written in multiple languages.
- **Non-executable modifications** – Some modifications are **not executable** (e.g., configuration changes, etc.)
- **Safety can be expensive** – what if the safe selection is still too **large** to run?

- **Ekstazi** – **safe** selection that people actually use: track which **files** each test touches at run time, and rerun a test only when one of its files changed. No coverage tool, no version-control hook – a JUnit plug-in.⁷
 - 32 projects, 615 revisions: end-to-end test time –32% on average, –54% for long suites.

⁷[ISSTA'15] M. Gligoric, L. Eloussi, D. Marinov. "Practical Regression Test Selection with Dynamic File Dependencies."

⁸[ICSE-SEIP'19] M. Machalica, A. Samylkin, M. Porth, S. Chandra. "Predictive Test Selection."

- **Ekstazi** – **safe** selection that people actually use: track which **files** each test touches at run time, and rerun a test only when one of its files changed. No coverage tool, no version-control hook – a JUnit plug-in.⁷
 - 32 projects, 615 revisions: end-to-end test time –32% on average, –54% for long suites.
- **Predictive test selection** (Facebook) – give up safety, and **learn** from history which tests a change is likely to **fail**.⁸
 - Testing cost **halved**; still reports $> 95\%$ of test failures and $> 99.9\%$ of faulty changes.
 - Not safe – but the **post-commit** run still catches the rest.

⁷[ISSTA'15] M. Gligoric, L. Eloussi, D. Marinov. "Practical Regression Test Selection with Dynamic File Dependencies."

⁸[ICSE-SEIP'19] M. Machalica, A. Samykin, M. Porth, S. Chandra. "Predictive Test Selection."

- Ekstazi rebuilds a test's **file** dependencies as it runs. A file is a coarse unit – touch any function in it and every test that imported it runs again.
- Going finer needs a **call graph**, and in **Python** there is none to trust: imports and types are decided at run time.

⁹[ISSTA'26] Y. Wang, M. Pradel, Z. Liu. "Names Are All You Need: Effective and Safe Regression Test Selection for Python."

- Ekstazi rebuilds a test's **file** dependencies as it runs. A file is a coarse unit – touch any function in it and every test that imported it runs again.
- Going finer needs a **call graph**, and in **Python** there is none to trust: imports and types are decided at run time.
- **NameRTS** – forget the call graph: link code elements by the **names** they use, and rerun a test if a changed element is **reachable** through those names.⁹
 - Names over-approximate, so two prunings cut them back using the **previous run** and the context of each name.
 - Skips **69.9%** of the test files – **2.5×** the file-level baseline – and is safe on **99.6%** of commits, against **76.6%**.

⁹[ISSTA'26] Y. Wang, M. Pradel, Z. Liu. "Names Are All You Need: Effective and Safe Regression Test Selection for Python."

- **Problem** – Regression test suite is **too large**

- **Problem** – Regression test suite is **too large**
- **Idea** – Execute tests following the order of their **importance**.

- **Problem** – Regression test suite is **too large**
- **Idea** – Execute tests following the order of their **importance**.
- **Solution** – **Prioritize** the test suite so that you get the most out of your regression testing whenever it gets stopped.

- **Problem** – Regression test suite is **too large**
- **Idea** – Execute tests following the order of their **importance**.
- **Solution** – **Prioritize** the test suite so that you get the most out of your regression testing whenever it gets stopped.
- Then, how to define the **importance** of each test case?

Without using any complicated techniques, we can simply prioritize the test using the following criteria:

- **Backwards** – Newer faults are more likely to be detected by newer tests.

Without using any complicated techniques, we can simply prioritize the test using the following criteria:

- **Backwards** – Newer faults are more likely to be detected by newer tests.
- **Random** – Randomly select the test cases without any bias.

Without using any complicated techniques, we can simply prioritize the test using the following criteria:

- **Backwards** – Newer faults are more likely to be detected by newer tests.
- **Random** – Randomly select the test cases without any bias.

If we want to prioritize test cases in a smarter way, what should we consider?

- Suppose we knew about all the faults in advance and we have the information about the **fault detection capability** of each test case. (Impossible but let's pretend)

- Suppose we knew about all the faults in advance and we have the information about the **fault detection capability** of each test case. (Impossible but let's pretend)
- Which test should we run first if we knew this? What next?

TC	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_9	f_{10}
t_1	✓				✓					
t_2	✓				✓	✓	✓			
t_3	✓	✓	✓	✓	✓	✓	✓			
t_4					✓					
t_5								✓	✓	✓

- Suppose we knew about all the faults in advance and we have the information about the **fault detection capability** of each test case. (Impossible but let's pretend)
- Which test should we run first if we knew this? What next?

TC	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_9	f_{10}
t_1	✓				✓					
t_2	✓				✓	✓	✓			
t_3	✓	✓	✓	✓	✓	✓	✓			
t_4					✓					
t_5								✓	✓	✓

- Obviously, we need to run t_3 first and then run t_5 as the next.

- Since we cannot know about the faults in advance, we can use a **surrogate** for the fault detection capability (e.g., coverage, etc.)

- Since we cannot know about the faults in advance, we can use a **surrogate** for the fault detection capability (e.g., coverage, etc.)
- Let's consider the **coverage** information as a **surrogate** for the fault detection capability.

TC	r_1	r_2	r_3	r_4	r_5	r_6	r_7	r_8	r_9	r_{10}
t_1	✓				✓					
t_2	✓				✓	✓	✓			
t_3	✓	✓	✓	✓	✓	✓	✓			
t_4					✓					
t_5								✓	✓	✓

- Since we cannot know about the faults in advance, we can use a **surrogate** for the fault detection capability (e.g., coverage, etc.)
- Let's consider the **coverage** information as a **surrogate** for the fault detection capability.

TC	r_1	r_2	r_3	r_4	r_5	r_6	r_7	r_8	r_9	r_{10}
t_1	✓				✓					
t_2	✓				✓	✓	✓			
t_3	✓	✓	✓	✓	✓	✓	✓			
t_4					✓					
t_5								✓	✓	✓

- We still need to run t_3 first and then run t_5 as the next.

- Since we cannot know about the faults in advance, we can use a **surrogate** for the fault detection capability (e.g., coverage, etc.)
- Let's consider the **coverage** information as a **surrogate** for the fault detection capability.

TC	r_1	r_2	r_3	r_4	r_5	r_6	r_7	r_8	r_9	r_{10}
t_1	✓				✓					
t_2	✓				✓	✓	✓			
t_3	✓	✓	✓	✓	✓	✓	✓			
t_4					✓					
t_5								✓	✓	✓

- We still need to run t_3 first and then run t_5 as the next.
- However, since the full coverage does **not guarantee** full fault detection, we still need to run remaining test cases after **resetting** the coverage information after running t_3 and t_5 .

- Since we cannot know about the faults in advance, we can use a **surrogate** for the fault detection capability (e.g., coverage, etc.)
- Let's consider the **coverage** information as a **surrogate** for the fault detection capability.

TC	r_1	r_2	r_3	r_4	r_5	r_6	r_7	r_8	r_9	r_{10}
t_1	✓				✓					
t_2	✓				✓	✓	✓			
t_3	✓	✓	✓	✓	✓	✓	✓			
t_4					✓					
t_5								✓	✓	✓

- We still need to run t_3 first and then run t_5 as the next.
- However, since the full coverage does **not guarantee** full fault detection, we still need to run remaining test cases after **resetting** the coverage information after running t_3 and t_5 .
- So, we need to run them in the order of t_2 , t_1 , and t_4 .

Average Percentage of Faults Detected (APFD)

- We can measure the effectiveness of the test case prioritization based on the **average percentage of faults detected** (APFD).¹⁰

¹⁰[TSE'01] G. Rothermel, R. H. Untch, C. Chu, M. J. Harrold. "Prioritizing Test Cases for Regression Testing."

Average Percentage of Faults Detected (APFD)

- We can measure the effectiveness of the test case prioritization based on the **average percentage of faults detected** (APFD).¹⁰
- Intuitively, APFD evaluates the effectiveness of the test case based on the **area** under the curve of the **fault detection rate**.


¹⁰[TSE'01] G. Rothermel, R. H. Untch, C. Chu, M. J. Harrold. "Prioritizing Test Cases for Regression Testing."

Average Percentage of Faults Detected (APFD)


- We can measure the effectiveness of the test case prioritization based on the **average percentage of faults detected (APFD)**.¹⁰
- Intuitively, APFD evaluates the effectiveness of the test case based on the **area** under the curve of the **fault detection rate**.

Example on
test suite and
faults exposed

Test Case	Fault							
	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8
t_A	•		•					•
t_B								
t_C		•		•	•			•
t_D	•	•				•		
t_E			•				•	



(a) APFD for test suite T_1




(b) APFD for test suite T_2

¹⁰[TSE'01] G. Rothmel, R. H. Untch, C. Chu, M. J. Harrold. "Prioritizing Test Cases for Regression Testing."

- **APFD_c** is a variant of APFD that weighs each test case by its **cost**: the x-axis is the test cost incurred, not the test count.¹¹

¹¹**[ICSE'01]** S. Elbaum, A. Malishevsky, G. Rothermel. "Incorporating Varying Test Costs and Fault Severities into Test Case Prioritization."

- **APFD_c** is a variant of APFD that weighs each test case by its **cost**: the x-axis is the test cost incurred, not the test count.¹¹




¹¹[ICSE'01] S. Elbaum, A. Malishevsky, G. Rothermel. "Incorporating Varying Test Costs and Fault Severities into Test Case Prioritization."

- A technique that **groups** objects such that objects in the same group are the most similar to each other.

- A technique that **groups** objects such that objects in the same group are the most similar to each other.
- Reduces the conceptual size of test suites by **clustering** test cases that are similar to each other.

- A technique that **groups** objects such that objects in the same group are the most similar to each other.
- Reduces the conceptual size of test suites by **clustering** test cases that are similar to each other.
- Provides insights into what is the **most common behavior** of the test cases.

- A technique that **groups** objects such that objects in the same group are the most similar to each other.
- Reduces the conceptual size of test suites by **clustering** test cases that are similar to each other.
- Provides insights into what is the **most common behavior** of the test cases.
- We can define a **distance** function between test cases based on diverse properties (e.g., coverage, etc.) and **visualize** the clustering results.



Cluster the test cases by their **coverage** (agglomerative hierarchical clustering), then let an expert rank the **clusters**, not the tests.¹²

Algorithm 1: Agglomerative Hierarchical Clustering

Input: A set of n test cases, T


Output: A dendrogram, D , representing the clusters

- (1) Form n clusters, each with one test case
- (2) $C \leftarrow \{\}$
- (3) Add clusters to C
- (4) Insert n clusters as leaf node into D
- (5) **while** there is more than one cluster
- (6) Find a pair of clusters with minimum distance
- (7) Merge the pair into a new cluster, c_{new}
- (8) Remove the pair of test cases from C
- (9) Add c_{new} to C
- (10) Insert c_{new} as a parent node of the pair into D
- (11) **return** D

¹²[ISSTA'09] S. Yoo, M. Harman, P. Tonella, A. Susi. "Clustering Test Cases to Achieve Effective and Scalable Prioritisation Incorporating Expert Knowledge."

Test Case Prioritization – Clustering

Cut the **dendrogram** at a height to get the clusters.¹³



¹³[ISSTA'09] S. Yoo, M. Harman, P. Tonella, A. Susi. "Clustering Test Cases to Achieve Effective and Scalable Prioritisation Incorporating Expert Knowledge."

Interleave: order the clusters, then take one test from each in turn.¹⁴

Algorithm 2: Interleaved Clusters Prioritisation

Input: An ordered set of k ordered clusters, OOC

Output: An ordered set of test cases, OTC

- (1) $OTC = \langle \rangle$
- (2) $i \leftarrow 1$
- (3) **while** OOC is not empty
- (4) Append $OOC_i(1)$ to OTC
- (5) Remove $OOC_i(1)$ from OOC_i
- (6) **if** OOC_i is empty **then** Remove OOC_i from OOC
- (7) $i \leftarrow (i + 1) \bmod k$
- (8) **return** OOC

$$OOC = [\quad \{A, B\}, \quad \{C, D, E\}, \quad \{F\} \quad]$$

¹⁴[\[ISSTA'09\]](#) S. Yoo, M. Harman, P. Tonella, A. Susi. "Clustering Test Cases to Achieve Effective and Scalable Prioritisation Incorporating Expert Knowledge."

Interleave: order the clusters, then take one test from each in turn.¹⁴

Algorithm 2: Interleaved Clusters Prioritisation

Input: An ordered set of k ordered clusters, OOC

Output: An ordered set of test cases, OTC


- (1) $OTC = \langle \rangle$
- (2) $i \leftarrow 1$
- (3) **while** OOC is not empty
- (4) Append $OOC_i(1)$ to OTC
- (5) Remove $OOC_i(1)$ from OOC_i
- (6) **if** OOC_i is empty **then** Remove OOC_i from OOC
- (7) $i \leftarrow (i + 1) \bmod k$
- (8) **return** OOC

$$OOC = [\quad \{A, B\}, \quad \{C, D, E\}, \quad \{F\} \quad]$$

$$OTC = [\quad A, C, F, B, D, E \quad]$$


¹⁴[\[ISSTA'09\]](#) S. Yoo, M. Harman, P. Tonella, A. Susi. "Clustering Test Cases to Achieve Effective and Scalable Prioritisation Incorporating Expert Knowledge."

FAST-R: a test is a vector of its **tokens**; sample the ones far apart – no coverage needed.¹⁵



¹⁵[ICSE'19] E. Cruciani, B. Miranda, R. Verdecchia, A. Bertolino. "Scalable Approaches for Test Suite Reduction."

FAST-R: a test is a vector of its **tokens**; sample the ones far apart – no coverage needed.¹⁵



Still a **syntactic** measure. Can we measure **semantic** similarity?

¹⁵[ICSE'19] E. Cruciani, B. Miranda, R. Verdecchia, A. Bertolino. "Scalable Approaches for Test Suite Reduction."

- Tokens are cheap but shallow – FAST-R buys its speed with **lost faults**.
- **LTM** – embed each test with a **pre-trained code model** (UniXcoder), take the **cosine** distance between embeddings, and search with a **genetic algorithm** for the most **dissimilar** subset.¹⁶
 - Still **black-box** – test code only, no coverage.
 - Detects **0.84** of the faults at a 50% budget, saves **41.7%** of the testing time on average, and runs **5×** faster than ATM, the previous similarity-based approach.
 - Similarity is now **learned**, not counted.

¹⁶[TSE'24] R. Pan, T. A. Ghaleb, L. C. Briand. "LTM: Scalable and Black-Box Similarity-Based Test Suite Minimization Based on Language Models."

- In **continuous integration (CI)** there is a new commit every few minutes and **no time** to collect coverage. What is left? The **history** of each test.

¹⁷[\[ISSTA'17\]](#) H. Spieker, A. Gotlieb, D. Marijan, M. Mossige. "*Reinforcement Learning for Automatic Test Case Prioritization and Selection in Continuous Integration.*"

- In **continuous integration (CI)** there is a new commit every few minutes and **no time** to collect coverage. What is left? The **history** of each test.
- **RETECS** – a **reinforcement learning** agent ranks each test from three numbers: its **duration**, when it **last ran**, and its recent **failures**.¹⁷
 - **Reward**: failures found early in this cycle. The agent learns from every CI run – no model of the code at all.
 - Three industrial datasets (ABB, Google): after a few hundred cycles it beats the fixed heuristics.

¹⁷[\[ISSTA'17\]](#) H. Spieker, A. Gotlieb, D. Marijan, M. Mossige. "Reinforcement Learning for Automatic Test Case Prioritization and Selection in Continuous Integration."

- In **continuous integration (CI)** there is a new commit every few minutes and **no time** to collect coverage. What is left? The **history** of each test.
- **RETECS** – a **reinforcement learning** agent ranks each test from three numbers: its **duration**, when it **last ran**, and its recent **failures**.¹⁷
 - **Reward**: failures found early in this cycle. The agent learns from every CI run – no model of the code at all.
 - Three industrial datasets (ABB, Google): after a few hundred cycles it beats the fixed heuristics.
- The surrogate is no longer coverage. It is **“this test failed recently”** – and that is what the next slide scales up to commits.

¹⁷[ISSTA'17] H. Spieker, A. Gotlieb, D. Marijan, M. Mossige. “*Reinforcement Learning for Automatic Test Case Prioritization and Selection in Continuous Integration.*”

- Most studies use CI builds from over **ten years ago**, with suites that run for **minutes**. **LRTS**: 21,255 builds from 10 large projects, average suite run **6.5 hours**.¹⁸

¹⁸[\[ISSTA'24\]](#) R. Cheng, S. Wang, R. Jabbarvand, D. Marinov. "Revisiting Test-Case Prioritization on Long-Running Test Suites."

- Most studies use CI builds from over **ten years ago**, with suites that run for **minutes**. **LRTS**: 21,255 builds from 10 large projects, average suite run **6.5 hours**.¹⁸
- 59 techniques re-evaluated. The winner: **run the fast tests that failed recently** – it beats every learning-based and text-retrieval-based technique.
- **Flaky** and always-failing tests blur the comparison, and the time-based order suffers the least from them.

¹⁸[\[ISSTA'24\]](#) R. Cheng, S. Wang, R. Jabbarvand, D. Marinov. "Revisiting Test-Case Prioritization on Long-Running Test Suites."



- A **DevOps** concept popularized by Google, more commonly and also known as: **Continuous Integration and Deployment (CI/CD)**



- A **DevOps** concept popularized by Google, more commonly and also known as: **Continuous Integration and Deployment (CI/CD)**
- Newest version of software is automatically deployed whenever **all tests pass**.



- A **DevOps** concept popularized by Google, more commonly and also known as: **Continuous Integration and Deployment (CI/CD)**
- Newest version of software is automatically deployed whenever **all tests pass**.
- Developers usually ensure that their commits are correct by executing test cases that are directly relevant at their **local machines**. This is sometimes called **pre-commit testing**.




- A **DevOps** concept popularized by Google, more commonly and also known as: **Continuous Integration and Deployment (CI/CD)**
- Newest version of software is automatically deployed whenever **all tests pass**.
- Developers usually ensure that their commits are correct by executing test cases that are directly relevant at their **local machines**. This is sometimes called **pre-commit testing**.
- Once changes are merged, the **CI system** automatically executes all relevant test cases, to ensure that individual changes correctly work with each other. This is sometimes called **post-commit testing**.

- All regression testing techniques are limited by the **cost** of running the test cases.

- All regression testing techniques are limited by the **cost** of running the test cases.
- In practice, we execute test cases in a **distributed** manner to minimize the cost of running the test cases.

- All regression testing techniques are limited by the **cost** of running the test cases.
- In practice, we execute test cases in a **distributed** manner to minimize the cost of running the test cases.
- Then, the testing time is only limited by the **slowest** test case.



- In very large development environments, the CI (Continuous Integration) pipeline is easily **flooded with commits**.

¹⁹[Velocity'11] J. Jenkins. "Velocity Culture." Amazon.com.

²⁰[ICSE-SEIP'17] A. Memon, Z. Gao, B. Nguyen, S. Dhanda, E. Nickell, R. Siemborski, J. Micco. "Taming Google-Scale Continuous Testing."

²¹[ICSE'18] J. Liang, S. Elbaum, G. Rothermel. "Redefining Prioritization: Continuous Prioritization for Continuous Integration."

- In very large development environments, the CI (Continuous Integration) pipeline is easily **flooded with commits**.
 - Amazon deployed to production every **11.6 seconds** on average, already in 2011.¹⁹

¹⁹[Velocity'11] J. Jenkins. "Velocity Culture." Amazon.com.

²⁰[ICSE-SEIP'17] A. Memon, Z. Gao, B. Nguyen, S. Dhanda, E. Nickell, R. Siemborski, J. Micco. "Taming Google-Scale Continuous Testing."

²¹[ICSE'18] J. Liang, S. Elbaum, G. Rothermel. "Redefining Prioritization: Continuous Prioritization for Continuous Integration."

- In very large development environments, the CI (Continuous Integration) pipeline is easily **flooded with commits**.
 - Amazon deployed to production every **11.6 seconds** on average, already in 2011.¹⁹
 - At Google, developers have waited **up to 9 hours** for the test results of a change.²⁰

¹⁹[Velocity'11] J. Jenkins. "Velocity Culture." Amazon.com.

²⁰[ICSE-SEIP'17] A. Memon, Z. Gao, B. Nguyen, S. Dhanda, E. Nickell, R. Siemborski, J. Micco. "Taming Google-Scale Continuous Testing."

²¹[ICSE'18] J. Liang, S. Elbaum, G. Rothmel. "Redefining Prioritization: Continuous Prioritization for Continuous Integration."

- In very large development environments, the CI (Continuous Integration) pipeline is easily **flooded with commits**.
 - Amazon deployed to production every **11.6 seconds** on average, already in 2011.¹⁹
 - At Google, developers have waited **up to 9 hours** for the test results of a change.²⁰
- **Prioritizing test cases** within one test suite makes little difference at this scale.

¹⁹[Velocity'11] J. Jenkins. "Velocity Culture." Amazon.com.

²⁰[ICSE-SEIP'17] A. Memon, Z. Gao, B. Nguyen, S. Dhanda, E. Nickell, R. Siemborski, J. Micco. "Taming Google-Scale Continuous Testing."

²¹[ICSE'18] J. Liang, S. Elbaum, G. Rothermel. "Redefining Prioritization: Continuous Prioritization for Continuous Integration."

- In very large development environments, the CI (Continuous Integration) pipeline is easily **flooded with commits**.
 - Amazon deployed to production every **11.6 seconds** on average, already in 2011.¹⁹
 - At Google, developers have waited **up to 9 hours** for the test results of a change.²⁰
- **Prioritizing test cases** within one test suite makes little difference at this scale.
- Instead, **prioritize the commits** to test.²¹

¹⁹[Velocity'11] J. Jenkins. "Velocity Culture." Amazon.com.

²⁰[ICSE-SEIP'17] A. Memon, Z. Gao, B. Nguyen, S. Dhanda, E. Nickell, R. Siemborski, J. Micco. "Taming Google-Scale Continuous Testing."

²¹[ICSE'18] J. Liang, S. Elbaum, G. Rothermel. "Redefining Prioritization: Continuous Prioritization for Continuous Integration."

- The technique is history-based prioritization: **commits** that touch code whose tests have **recently failed** are tested first. If tests really fail, this ensures **quicker feedback** to the responsible developers.

- The technique is history-based prioritization: **commits** that touch code whose tests have **recently failed** are tested first. If tests really fail, this ensures **quicker feedback** to the responsible developers.
- Assumptions
 - Commits are **independent** from each other. In high volume environment, this is not unrealistic.
 - **Relationships** between code and test cases are known in advance (i.e., which test covers which parts of the code).

- **Minimization:** Keyword is **redundancy** – **Minimize** test suite by removing redundant test cases according to the definition of redundancy.
- **Selection:** Keyword is **safety** – **Select** test cases that are related to the recent changes conservatively to avoid possible regression faults.
- **Prioritization:** Keyword is **surrogate** – **Prioritize** test cases based on the surrogate for the fault detection capability to maximize the fault detection rate early.

- Mutation Testing

Jihyeok Park

jihyeok_park@korea.ac.kr

<https://plrg.korea.ac.kr>